

University of Colorado at Colorado Springs

Spring 2021 Operating Systems Project 1

System Calls and Processes

Instructor: Yanyan Zhuang
Total Points: 100
Out: 2/1/2021
Due: 11:59 pm, Monday, 2/22/2021

Introduction

The purpose of this project is to study how system calls and processes are managed by the Linux OS. The objectives of this project is to:

1. Get familiar with the Linux environment.
2. Learn how to build a Linux kernel.
3. Learn how to add a new system call in the Linux kernel.
4. Learn how to obtain various information for a running process from the Linux kernel.

Project submission

For each project, please create a zipped file containing the following items, and submit it to Canvas.

1. A report that includes (1) the (printed) full names of the project members, and the statement: **We have neither given nor received unauthorized assistance on this work**; (2) the name of your virtual machine (VM), and the password for the **instructor** account (details see *Accessing VM* below); and (3) a brief description about how you solved the problems and what you learned. For example, you may want to make a typescript of the steps in building the Linux kernel, or the changes you made in the Linux kernel source for adding a new system call, or screenshots of your result. The report can be in txt, doc, or pdf format. Please use the template provided in Canvas Modules → Tutorial and Template → writeup-template.docx (thanks to Brandon who created it), or download the template at http://cs.uccs.edu/~yzhuang/CS4500_5500/writeup-template.docx
2. The source code including both the kernel-level code you added and the user-level testing program.¹

¹In case there are issues accessing your VM, submitting your source code separately would allow the instructor to give partial credit.

Accessing VM. Each team should specify the name of your VM that the instructor can login to check the project results. Please create a new user account called `instructor` in your VM, and place your code in the home directory of the instructor account (i.e., `/home/instructor`). Make sure the `instructor` has the appropriate access permission to your code. In your project report, include your password for the `instructor` account.

Tasks

Task 0.0: Use the provided template for writeup (5 points)

Task 0: Building the Linux kernel (20 points)

The version of the CentOS kernel we use is quite old. However, this version is small, and therefore takes less time to recompile. To mitigate the issue of an old kernel, we need to do a few things (you only need to do them once).

First, edit `/etc/yum.repos.d/CentOS-Base.repo` and replace all `mirrorlist` as root (note that the `#` prompt indicates that the command requires root privilege. `$` indicates normal user privilege):

```
# cd /etc/yum.repos.d/
# cp CentOS-Base.repo CentOS-Base.repo.old
# vi CentOS-Base.repo
```

Replace sections:

```
[base]
[updates]
[extras]
```

With the following (the best way is to use `ssh` and copy the following, **not through the datacenter web interface**):

```
[base]
name=CentOS-$releasever - Base
baseurl=https://vault.centos.org/6.10/os/$basearch/
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-CentOS-6

# released updates
[updates]
name=CentOS-$releasever - Updates
baseurl=https://vault.centos.org/6.10/updates/$basearch/
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-CentOS-6

# additional packages that may be useful
[extras]
name=CentOS-$releasever - Extras
baseurl=https://vault.centos.org/6.10/extras/$basearch/
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-CentOS-6
```

Then clean yum cache:

```
# yum clean all
```

The detailed steps are found here:

<https://arstech.net/centos-6-error-yumrepo-error-all-mirror-urls-are-not-using-ftp-http/>

Step 1: Getting the Linux kernel code

Before you download and compile the Linux kernel source, make sure that you have development tools installed on your system.

In CentOS, install these software using yum:

```
# yum install -y gcc ncurses-devel make wget
```

In the past some VMs don't have perl pre-installed. If that's the case, do `yum install perl` with root privilege.

Visit <http://kernel.org> and download the source code of your current running kernel. To obtain the version of your current kernel, type:

```
$ uname -r
```

```
$ 2.6.32-220.el6.i686
```

Then, download kernel 2.6.32 and extract the source:

```
$ wget http://www.kernel.org/pub/linux/kernel/v2.6/linux-2.6.32.tar.gz  
--no-check-certificate
```

Note the above is one single command.

```
$ tar xvzf linux-2.6.32.tar.gz
```

We will use `LINUX_SOURCE` to refer to the top directory of the kernel source, i.e., `linux-2.6.32/`.

Step 2: Configure your new kernel

It's a good idea to take a snapshot of your VM in case things get messed up. However, please keep the number of snapshots under 4.

Before compiling the new kernel, a `.config` file needs to be generated in the top directory of the kernel source. To generate the config file, type this **in `LINUX_SOURCE`**:

```
$ make localmodconfig
```

Make sure that you hit "enter" when prompted for questions.

Step 3: Compile the kernel

In `LINUX_SOURCE`, compile to create a compressed kernel image:

```
$ make -j8
```

Replace 8 in `make -j8` with the number of cores you have for the VM (typically a number between 4-8. if you don't know, this wouldn't affect you, just an optimization option). If you see errors from `make`, **DO NOT ignore the errors**. Read the error message and try to trouble shoot.

Step 4: Install the kernel

Compile and install kernel modules (to become the root user, use the `su` command):

```
$ su
```

```
# make modules
```

```
# make modules_install
```

Again, in the above, `$` means you are a normal user, and `#` means you are root. Now install the kernel:

```
# export KERNELVERSION=$(make kernelversion)
# cp arch/x86/boot/bzImage /boot/vmlinuz-$KERNELVERSION
# cp System.map /boot/System.map-$KERNELVERSION
```

When you run the two `cp` commands above, and if you are asked whether to overwrite the existing files, make sure that you hit `y` to overwrite. The first time you ran it, you wouldn't see this message.

```
# new-kernel-pkg --mkinitrd --install $KERNELVERSION
```

The kernel image and other related files have been installed into the `/boot` directory.

Step 5: Reboot your VM

Reboot to the new kernel:

```
# reboot
```

Please make sure that you select the new kernel that has been built (i.e., 2.6.32).² **Note that if you use ssh, a reboot would disconnect your laptop with the VM. So you may use the GUI interface to reboot.** After reboot, always check if you have the new kernel:

```
$ uname -r
$ 2.6.32
```

Task 1: Add a new system call to the Linux kernel (20 points)

In this task, we will add a simple system call `helloworld` to the Linux kernel. The system call prints out a `Hello World!` message to the `syslog`. You need to implement the system call in the kernel and write a user-level program to test your new system call.

Step 1: Implement your system call

Change your current working directory to the kernel source directory.

```
$ cd LINUX_SOURCE
```

Make a new directory called `my_source` to contain your implementation:

```
$ mkdir my_source
```

Create a C file and implement your system call here:

```
$ touch my_source/sys_helloworld.c
```

Edit the source code to include the following implementation (be careful about the quotation marks – type the `printk` function by hand rather than copy&paste to avoid encoding issues; also note that the quotations in `printk` are double quotes rather than two single quotes):

```
$ vim my_source/sys_helloworld.c
#include <linux/kernel.h>
#include <linux/sched.h>
asmlinkage int sys_helloworld (void)
{
```

²In the past, some people had the issue of not having an option to select a new build upon reboot. If this happens, try holding down the shift key while reboot. **When you see a message “Boot CentOS (2.6.32-220.el6.i686)” hit enter right away.**

```

    printk(KERN_EMERG "Hello World!\n");
    return 0;
}

```

Add a Makefile to the `my_source` folder:

```
$ touch my_source/Makefile
```

```
$ vim my_source/Makefile
```

```

#
#Makefile of the new system call
#
obj-y := sys_helloworld.o

```

Modify the Makefile in the top directory `LINUX_SOURCE` to include the new system call in the kernel compilation:

```
$ vim Makefile
```

Find the line where `core-y` is defined and add the `my_source` directory to it:

```
core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/ my_source/
```

Add a function pointer to the new system call in `arch/x86/kernel/syscall_table_32.S`:

```
$ vim arch/x86/kernel/syscall_table_32.S
```

Add the following line to the end of this file:

```
.long sys_helloworld /* 337 */
```

Now you have defined `helloworld` as system call 337.

Register your system call with the kernel by editing `arch/x86/include/asm/unistd_32.h`. If the last system call in `arch/x86/kernel/syscall_table_32.S` has an id number 336, then our new system call should be numbered as 337. Add this to the file `unistd_32.h`, right below the definition of system call 336:

```
#define __NR_helloworld 337
```

Modify the total number of system calls to $337 + 1 = 338$:

```
#define NR_syscalls 338
```

Declare the system call routine. Add the following to the end of `arch/x86/include/asm/syscall.h` (right before the line `CONFIG_X86_32`, near the end of the file):

```
asmlinkage int sys_helloworld(void);
```

Repeat steps 3, 4 and 5 in task 0 to re-compile, install and reboot to the new kernel.

Step 2: Write a user-level program to test your system call

Go to your home directory and create a test program `test_syscall.c`:

```

#include <linux/unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <stdio.h>
#define __NR_helloworld 337
int main(int argc, char *argv[]) {

```

```

    syscall(__NR_helloworld);
    return 0;
}

```

Compile the program:

```
$ gcc test_syscall.c -o test_syscall
```

Test the new system call by running:

```
$ ./test_syscall
```

If your program hangs, hit **enter** to continue. The test program will call the new system call and output a helloworld message at the tail of the output of `dmesg` (try `dmesg | tail` to see the output).

Task 2: Extend your new system call to print the calling process's information (25 points)

Follow the instructions we discussed above and implement another system call `print_self`. This system call identifies the current process at the user-level and print out various information of the process.

Implement the `print_self` system call and print out the following information of the current process:

- Process id and program name
- The same above information of its parent processes until `init`

HINT: The macro `current` returns a pointer to the `task_struct` of the current running process. See the following link: <http://linuxgazette.net/133/saha.html> Please pay attention to the last example that traces from the current process back to `init`.

Task 3: Extend your new system call to print the information of an arbitrary process identified by its PID (30 points)

Implement another system call `print_other` to print the information for an arbitrary process. The system call takes a process pid as its argument and outputs the above information (same as Task 2) of this process.

HINT: You can start from the `init` process and iterate over all the processes. For each process, compare its pid with the target pid. If there is a match, return the pointer to this `task_struct`. A better approach would be to use the functions that Linux provides to look up the process in the process table. You may refer to the lecture slides.

To get a valid PID for a process that's active, you may use command `pgrep bash`. `bash` is the default command line shell. `pgrep bash` will provide PIDs of `bash`. You can choose any of them if you get more than one. More information about command line tools to get process information:

<http://linuxcommando.blogspot.com/2007/11/how-to-look-up-process.html>

In C, command line arguments passed in will be treated as strings. You need to convert them to other types that you need, like integers.